

GEONSIK MOON

(917) 257-7860 | geonsik.moon@columbia.edu | linkedin.com/in/gsmoon97 | github.com/gsmoon97

EDUCATION

Columbia University <i>Master of Science, Computer Science</i>	Expected Dec 2026 New York, NY
National University of Singapore <i>Bachelor of Computing (Honors), Computer Science</i>	Jun 2022 Singapore

• Machine Learning Track
• IBM Research Project Collaboration on Granite Speech Model Optimization

• Awarded the Certificate of Distinction in the Artificial Intelligence Focus Area
• Awarded the Certificate of Merit in the Database Systems Focus Area

EXPERIENCE

LLM Training Operation Specialist <i>ByteDance</i>	May 2024 – Aug 2025 Singapore
• Diagnosed systemic failure patterns in software engineering agents through analysis of 3K+ evaluation outputs, producing root-cause reports and actionable patch roadmaps that improved Supervised Fine-Tuning (SFT) performance by 30% to achieve parity with state-of-the-art models • Directed end-to-end data operations for a multi-agent LLM system by supervising ~70 annotators, implementing calibration protocols and cross-team coordination, successfully delivering 17K+ high-quality data points for Reinforcement Learning (RL) training that contributed to development of CodeContests+ benchmark • Led collaboration with subject matter experts for AetherCode benchmark by curating evaluation datasets from top coding competitions (IOI, ICPC) with comprehensive test suites, establishing a rigorous new standard in code reasoning where top models achieve only 35.5% Pass@1 • Managed SFT project for instruction-tuning, curating 2K+ expert-validated data points to align foundational models for complex coding capabilities across JavaScript, Python, SQL, and Go • Developed optimization tools in Python and implemented prompt engineering techniques (few-shot, Chain-of-Thought (CoT)) to streamline training workflows for data format conversion and trajectory analysis	
NLP Research Assistant <i>National University of Singapore</i>	Sep 2022 – Feb 2024 Singapore

• Co-authored two *ACL 2024* publications, advancing the state-of-the-art of LLM capabilities for downstream tasks of timeline summarization and understanding word semantics
• Implemented end-to-end LLM pipelines utilizing LangChain and Chroma, seamlessly integrating open-source models for incremental clustering algorithm deploying LLM-based pairwise classification
• Fine-tuned LLMs (Mistral, Llama 2, FLAN-T5) using LoRA-based PEFT with 4-bit quantization, employing various hyperparameter tuning based on performance tracking from Weights & Biases
• Published two system demonstration papers on GEC web applications in *EACL 2023* and *IJCNLP-AACL 2023*, detailing innovative design strategies to boost NLP pipeline efficiency and accuracy
• Engineered scalable, production-grade web applications for Grammatical Error Correction (GEC) by leveraging containerized microservices (Docker) and modern web frameworks (Flask, Bootstrap)

Machine Learning Engineer
Apple (via TransPerfect)

• Served as a vendor language engineer for the Global Siri Team, performing comprehensive error analyses and fine-tuning model inputs to enhance the model's natural language understanding
• Explored synthetic data augmentation via transformer-driven methods (BERT, T5), utilizing back-translation and context-aware paraphrasing to expand linguistic coverage while preserving semantic integrity
• Automated a dialogue optimization pipeline using scheduled Python scripts and regex pattern matching, reducing extraneous utterances by ~30% and enhancing response speed and relevance

TECHNICAL SKILLS

Programming Languages: Python, Java, JavaScript/TypeScript, C/C++, Go, SQL

ML Frameworks: PyTorch, TensorFlow, Hugging Face Transformers, scikit-learn

LLM Training & Evaluation: LoRA/PEFT, RLHF/SFT, LM Eval Harness, Weights & Biases

LLM Infrastructure: LangChain/LangGraph, vLLM, Ollama, llama.cpp, Chroma, Pinecone

Cloud & DevOps: AWS (Bedrock, EC2, S3), GCP, Docker, CUDA, Git